

AIQG — Investor Positioning Analysis

Status: Working analysis — 2026-06-12. **Inputs:** [AIQG-EXTENSION.md](#), [AIQG-AGENT-FLOW-ATTRIBUTION.md](#), [AIQG-AGENT-IDENTITY-RESEARCH.md](#), [AIQG-EXPERIMENTS-RUNNER.md](#).
Companion: [AIQG-PATENT-ANALYSIS.md](#)

Summary

The pitch is not any one feature — it's that the four pieces (extension / attribution / identity / experiments) close a loop nobody else closes: **attribute** → **score** → **experiment** → **optimize**, **all at the gateway, with zero client code changes**.

The wedge: agentless attribution in an instrumented market

The LLM gateway/observability category is crowded (Helicone, Langfuse, LangSmith, Portkey, Kong/Cloudflare AI gateways) — but every competitor shares one assumption AIQG breaks: **the customer must instrument their code** (headers, SDKs, trace context). The prior-art research confirms this is universal (“today's attribution is self-asserted everywhere”).

AIQG's inferred tier means a customer points DNS at the gateway and gets per-agent cost/quality attribution on day one with zero code changes. This is the agentless-vs-agent-based monitoring story replayed for AI — and historically the agentless entrant wins the land-grab because it removes the adoption tax.

Three investor-legible storylines

1. FinOps for AI agents

Per-agent and per-flow cost rollup, plus the avoidable-cost decomposition (direct / induced / genuine waste) from the CLEAR extension. *“Which of our 40 agents is burning the budget, and how much of that spend is avoidable”* is a CFO question nobody can currently answer without an instrumentation project.

The agent registry (“14 distinct agents detected, 11 unnamed”) doubles as **shadow-AI discovery**, which sells to the CISO in the same meeting — and the asserted-vs-inferred impersonation cross-check gives security a reason to keep the product after the cost story lands it.

2. The closed loop (measurement → action)

Competitors stop at dashboards. The Experiments Runner makes the measurement actionable: detect that a workflow is over-provisioned on gpt-4o → run a guardrailed, sticky, non-inferiority A/B against 4o-mini → quantified verdict. That's Statsig/LaunchDarkly-for-LLM-routing, executed at the gateway so it again needs no app changes.

The **Gatekeeper-impact experiment** is a unique proof-point: we can *show customers data* that the security layer doesn't degrade answer quality — a claim every competing AI firewall asserts and none can demonstrate.

3. Positioned for the authenticated-agent-identity wave

The schema ships self-asserted (matching the whole market) but reserves `agent_identity_verified` + credential refs so AIP/IBCT, Google A2A, or token-bound identity slots in without migration. Agent identity is an emerging 2026 standards fight; being the gateway that already records

graded identity confidence is a cheap option on that future. OTel gen_ai.* alignment also de-risks the “proprietary schema lock-in” objection.

Moat shape

- **Data network effect:** the fingerprint registry compounds — every labeled cluster and every framework signature (CrewAI / LangChain / AutoGen idioms) improves inference for all tenants; competitors can’t shortcut it.
- **Deterministic floor:** the token-echo and prefix-chain tiers are proof-not-probability, so the headline capability holds even if fingerprint clustering underperforms.
- **IP option:** a provisional on the inference tier (see AIQG-PATENT-ANALYSIS.md) turns the wedge into a defensibility story beyond “we built a nice proxy.”

Risks an investor will probe (be ready)

Risk	Honest counter
FTO overhang on metering/billing (4 located USPTO patents, claims unread)	Scoped: gates billing use only, not observability; FTO pass is planned and cheap relative to the raise
Inference precision/recall unproven	The validation harness with known ground truth (cmd/demo-traffic --untagged) is already built; deterministic tiers set a high floor
OTel GenAI conventions are experimental	Names pinned; additive schema means renames are migrations, not breakage
Hyperscalers bundle basic gateway attribution	They’ll ship header-based (self-asserted) attribution; the inferred tier + registry is the differentiated layer
Crowded category	The wedge is adoption cost, not feature count — zero-code-change attribution is structurally hard for SDK-first incumbents to retrofit

Concrete next steps

1. **Run the untagged validation** and publish precision/recall per inference tier internally — the single best artifact for the pitch deck.
2. **FTO claim-read** before diligence surfaces it for us.
3. Keep the inferred-attribution design non-public until the provisional-vs-defensive-publication decision is made.