

AIQG — Investor Positioning Analysis

Status: Working analysis — 2026-06-12; **IP section revised 2026-06-19** to reflect the completed FTO claim-read and the six-candidate provisional brief. **Inputs:** [AIQG-EXTENSION.md](#), [AIQG-AGENT-FLOW-ATTRIBUTION.md](#), [AIQG-AGENT-IDENTITY-RESEARCH.md](#), [AIQG-EXPERIMENTS-RUNNER.md](#). **Companions:** [AIQG-PATENT-ANALYSIS.md](#), [AIQG-PROVISIONAL-BRIEF.md](#) (Candidates 1–6), [AIQG-FTO-CLAIM-READ.md](#), [AIQG-PATENT-SCOUT-AUTOLEARNING.md](#), [AIQG-PATENT-ML-IN-LOOP.md](#)

Summary

The pitch is not any one feature — it’s that the four pieces (extension / attribution / identity / experiments) close a loop nobody else closes: **attribute** → **score** → **experiment** → **optimize**, **all at the gateway, with zero client code changes**.

The wedge: agentless attribution in an instrumented market

The LLM gateway/observability category is crowded (Helicone, Langfuse, LangSmith, Portkey, Kong/Cloudflare AI gateways) — but every competitor shares one assumption AIQG breaks: **the customer must instrument their code** (headers, SDKs, trace context). The prior-art research confirms this is universal (“today’s attribution is self-asserted everywhere”).

AIQG’s inferred tier means a customer points DNS at the gateway and gets per-agent cost/quality attribution on day one with zero code changes. This is the agentless-vs-agent-based monitoring story replayed for AI — and historically the agentless entrant wins the land-grab because it removes the adoption tax.

Three investor-legible storylines

1. FinOps for AI agents

Per-agent and per-flow cost rollup, plus the avoidable-cost decomposition (direct / induced / genuine waste) from the CLEAR extension. *“Which of our 40 agents is burning the budget, and how much of that spend is avoidable”* is a CFO question nobody can currently answer without an instrumentation project.

The agent registry (“14 distinct agents detected, 11 unnamed”) doubles as **shadow-AI discovery**, which sells to the CISO in the same meeting — and the asserted-vs-inferred impersonation cross-check gives security a reason to keep the product after the cost story lands it.

2. The closed loop (measurement → action)

Competitors stop at dashboards. The Experiments Runner makes the measurement actionable: detect that a workflow is over-provisioned on gpt-4o → run a guardrailed, sticky, non-inferiority A/B against 4o-mini → quantified verdict. That’s Statsig/LaunchDarkly-for-LLM-routing, executed at the gateway so it again needs no app changes. **Scout** sharpens this: it mines the customer’s own traffic for the cheaper-model swaps worth testing and estimates the **time-to-verdict before launch** (“this swap would take ≈3 weeks to prove at 5% exposure”), so the optimization loop opens with a prioritized, feasibility-sized list rather than a blank experiment form.

The **Gatekeeper-impact experiment** is a unique proof-point: we can *show customers data* that the security layer doesn’t degrade answer quality — a claim every competing AI firewall asserts and none can demonstrate.

3. Positioned for the authenticated-agent-identity wave

The schema ships self-asserted (matching the whole market) but reserves `agent_identity_verified` + credential refs so AIP/IBCT, Google A2A, or token-bound identity slots in without migration. Agent identity is an emerging 2026 standards fight; being the gateway that already records graded identity confidence is a cheap option on that future. OTel `gen_ai.*` alignment also de-risks the “proprietary schema lock-in” objection.

Moat shape

- **Data network effect:** the fingerprint registry compounds — every labeled cluster and every framework signature (CrewAI / LangChain / AutoGen idioms) improves inference for all tenants; competitors can’t shortcut it.
- **Deterministic floor:** the token-echo and prefix-chain tiers are proof-not-probability, so the headline capability holds even if fingerprint clustering underperforms.
- **IP option, now de-risked:** the defensibility story is no longer hypothetical. A **six-candidate provisional** is drafted (AIQG-PROVISIONAL-BRIEF.md) — the zero-cooperation attribution tier (content-propagation graph, tool-call echo, graded identity ladder, fingerprint registry) plus identity-ladder-keyed sticky experimentation and Scout’s pre-launch time-to-verdict — and the **FTO claim-read is done** (AIQG-FTO-CLAIM-READ.md): no high-risk read against the four located patents, and the strongest candidate (content-propagation graph) is clear of all four. A second filing track (autolearning + SLM/ML-in-loop, e.g. self-labeling a classifier from the protocol’s forced echo) is scoped behind a bandit-routing prior-art search.

Risks an investor will probe (be ready)

Risk	Honest counter
FTO overhang on metering/billing (4 located USPTO patents)	Claim-read now done (AIQG-FTO-CLAIM-READ.md): no high-risk read on any of our six techniques. Only Dynatrace US 10,924,326 is a medium on per-agent metering — gates <i>billing-grade</i> use, not observability, and our inference-only (no injected correlation id) design is the documented design-around. Counsel opinion on that one patent is the remaining cheap step
Inference precision/recall unproven	The validation harness with known ground truth (<code>cmd/demo-traffic --untagged</code>) is already built; deterministic tiers set a high floor
OTel GenAI conventions are experimental	Names pinned; additive schema means renames are migrations, not breakage
Hyperscalers bundle basic gateway attribution	They’ll ship header-based (self-asserted) attribution; the inferred tier + registry is the differentiated layer
Crowded category	The wedge is adoption cost, not feature count — zero-code-change attribution is structurally hard for SDK-first incumbents to retrofit

Concrete next steps

1. **Run the untagged validation** and publish precision/recall per inference tier internally — the single best artifact for the pitch deck, and the reduction-to-practice exhibit for the provisional.
2. **File the six-candidate provisional** (AIQG-PROVISIONAL-BRIEF.md) before any public disclosure; commission the counsel opinion on Dynatrace 10,924,326 ahead of any billing-grade metering. (*FTO claim-read — done.*)
3. **Prototype one ML-in-loop claim** for the second track — the self-labeling classifier (M1) or feedback-recalibrated SLM judge (M2) — to turn the autolearning story from idea into a demonstrable, defensible asset.
4. Keep the inferred-attribution, autolearning, and ML-in-loop designs non-public until each provisional is on file.